


Using Parallelism to Speed Up Cluster Validation

The Purpose of Clustering:

- “...clustering analyzes data objects without consulting class labels.” Data Mining: Concepts and Techniques.
- Associates unlabeled (unclassified) data into groups based on some distance measure.
- Distance, proximity, close-ness and similarity all measure roughly the same thing and there are many ways to measure distance.
- There are *many* decisions to make when clustering data.

Cluster Validity:



6-Means clustering of 100,000 points generated around 10 random centers.

- How can we tell that a clustering is “good”?
- Judgement is easy in low dimensional data but much harder when visualization is more difficult.
- Validation methods can either be external (which rely on external knowledge) or internal (which only relies on the measures from the data)

Dunn Index:

- First proposed by JC Dunn in the Journal of Cybernetics in 1974.
- Indicates the ratio between the minimum distance between any two clusters and the the maximum of the mean distance between points in any cluster.
- **High value** indicates well separated clusters.
- Because it relies on a “max value in the denominator” it is very sensitive to one bad cluster.
- $\delta(C_i, C_j)$ is the inter-cluster distance measure.
 - Center to center for this presentation.
- Δ_k is the max distance intra-cluster distance.
 - Point to center for this presentation.

$$DI_m = \frac{\min_{1 \leq i < j \leq m} \delta(C_i, C_j)}{\max_{1 \leq k \leq m} \Delta_k}$$

Davies-Bouldin Index:

- Defined by David Davies and Donald Bouldin in 1979 in *IEEE Transactions on Pattern Analysis and Machine Intelligence*.
- Represents a ratio between distances between points to centroids and centroids to centroids.
- **Low values** indicate strongly dissimilar clusters.
- S_i is the mean of all distances from a point to its own centroid i .
- $M_{i,j}$ is the distance between cluster i and j .
- N is the number of clusters.



$$R_{i,j} = \frac{S_i + S_j}{M_{i,j}}$$

$$D_i \equiv \max_{j \neq i} R_{i,j}$$

$$DB \equiv \frac{1}{N} \sum_{i=1}^N D_i$$

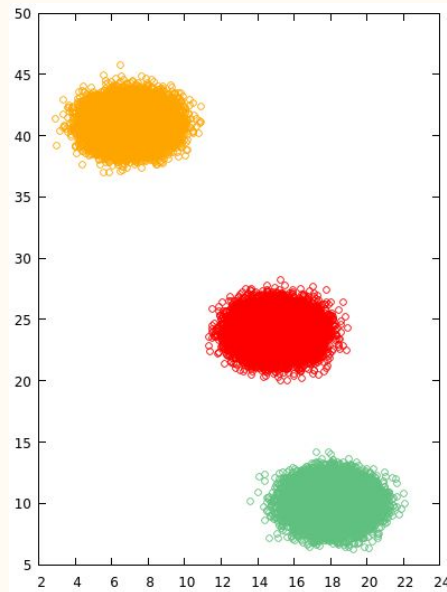

Parallelizing These Approaches:

- Implement a Dunn Index and DB Index calculator in C++.
- Seperate loops to minimize dependencies.
- Introduce OpenMP
- Merge parallel executions or protect memory assets as necessary to maintain data integrity.
- Insert time checking code.
- All code was compiled with gcc at at -O0.





Test Data Set 1:

- 100,000 points
- 2 dimensions
- Generated in gaussian distributions around 3 centers.
- 2 test clustering:
 - $k=2$
 - $k=3$


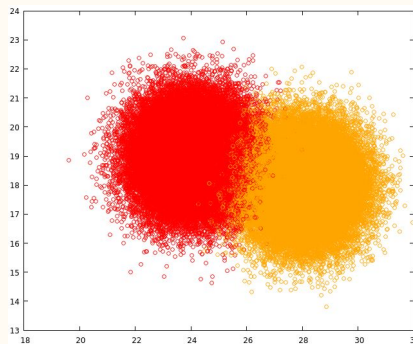


Test Data Set 2:

- 100,000 points
- 10 dimensions
- Generated in gaussian distributions around 2 centers.
- 2 test clustering:
 - $k=2$
 - $k=3$



Projected onto dimension 1 and 2




Projected onto dimension 2 and 3


Test Data Set 2:

- 100,000 points
- 10 dimensions
- Generated in gaussian distributions around 10 centers.
- 3 test clustering:
 - $k=8$
 - $k=10$
 - $k=12$


8 clusters






10 clusters



12 clusters



Projected onto dimensions 1 and 2




Projected onto dimensions 2 and 3

Set 1 Results:

	2 clusters	3 clusters
Dunn Index	2.287	2.988
DB Index	0.164	0.084




Best
Fit



Set 2 Results:

	2 clusters	3 clusters
Dunn Index	3.668	0.239
DB Index	0.124	1.308




Set 3 Results:

	8 clusters	10 clusters	12 clusters
Dunn Index	0.021	0.025	0.027
DB Index	0.674	1.5	1.327




Best
Fit




Dunn Index Performance:

Run Time of Dunn Index on Set 1




Run Time of Dunn Index on Set 2




Davies-Bouldin Index Performance:

Run Time of DB Index on Set 1



Run Time of DB Index on Set 2




What's Wrong with DB?

- Outer loop iterating across all clusters means the benefit of parallelism is lost.

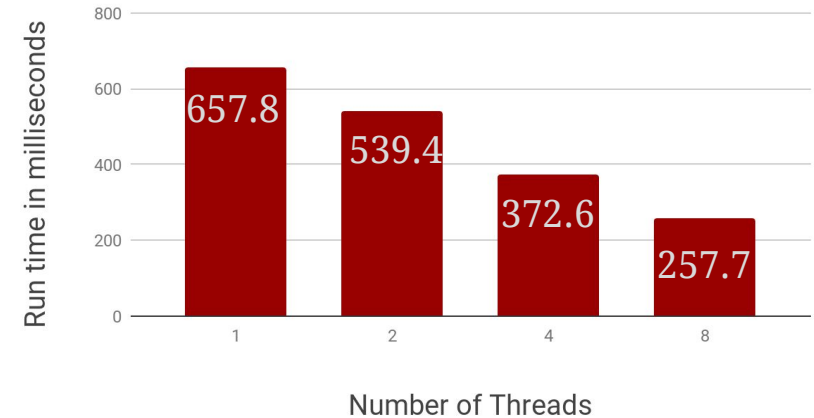
```
#pragma omp parallel for
for (i = 0; i < numClusters; ++i){ //divide sum by points to get centers
    centers[i].values = new float[dimensions];
    int j;
    for (j = 0; j < dimensions; ++j){
        centers[i].values[j] = (sums[i][j] / totalPoints[i]);
    }
    centers[i].cluster = i+1;
}
```

Set 3 Results:

Mean Run Time of Dunn Index on Set 3



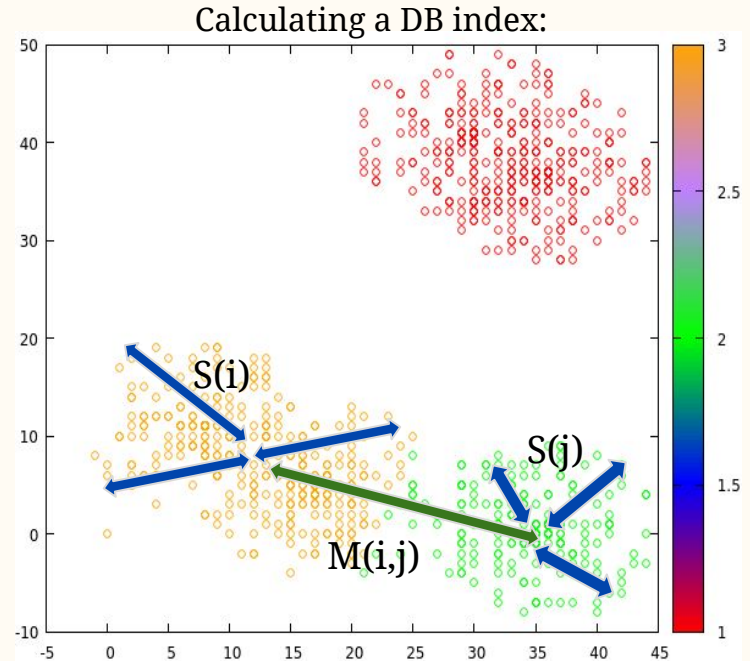
Mean Run Times of DB Index on Set 3



Run times averaged across three runs for three values of k display the expected increase in performance.

Applications:

- Performance improvement will be multiplied in hierarchical systems and the results will be more meaningful.
- Non-deterministic clustering algorithms could have several runs evaluated quickly and easily.
- Selecting a clustering algorithm is often a very difficult problem. More powerful and convenient tools to evaluate clusters makes the problem easier.




Future Work and Refinements:

- Re-run tests on distinct architectures.
- Re-run tests with LLVM.
- Continue to refactor code to derive more benefit from parallelism.
Explore OMP reduction further.



Photograph by Rama, Wikimedia Commons,
Cc-by-sa-2.0-fr, CC BY-SA 2.0 fr,
<https://commons.wikimedia.org/w/index.php?curid=14619982>



Dilbert.com DilbertCartoonist@gmail.com



11-24-14 © 2014 Scott Adams, Inc. /Dist. by Universal Uclick



Thank you for your attention.

Questions or Comments?

Code and data available at: <https://github.com/gentry-atkinson/cs7331.git>